

ADAPTIVE CO-SCHEDULER FOR HIGHLY DYNAMIC RESOURCE

by

Kartik Vedalaveni

A THESIS

Presented to the Faculty of

The Graduate College at the University Of Nebraska

In Partial Fulfilment of Requirements

For the Degree of Master Of Science

Major: Computer Science

Under the Supervision of Dr. David Swanson

Lincoln, Nebraska

May, 2013

# ADAPTIVE CO-SCHEDULER FOR HIGHLY DYNAMIC RESOURCE

Kartik Vedalaveni, M.S

University Of Nebraska, 2013

Adviser: Dr. David Swanson

Schedulers come with plethora of features and options for customization that fulfills myriad goals of clusters and data centers . Most state of the art schedulers do not take into account load of resources like RAM, I/O or Network for scheduling purpose. Often there is a need to extend these schedulers to solve situations arising from these new use cases either by writing a plugin or by Co-Scheduling. One such issue is when any resources like I/O, RAM or Network is throttled and the degradation that occurs as a result of it . With increase in number of entities concurrently using the resource, there is a need to monitor and schedule concurrent and unmanaged access to any given resource to prevent degradation. These issues that we encounter in real life at Holland Computing Center (CITE) are the basis and motivation for tackling this problem and develop an adaptive approach for scheduling that is aware of multi-resource throttling, load-balancing across multiple sites and degradation with a goal to run clusters at high efficiency and share resources fluidly.

## DEDICATION

Dedicated to

## ACKNOWLEDGMENTS

Thanks

# Contents

<b>Contents</b>	<b>v</b>
<b>List of Figures</b>	<b>vii</b>
<b>List of Tables</b>	<b>viii</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Co-Scheduler . . . . .	4
<b>2 Background</b>	<b>5</b>
2.1 High-Throughput Computing . . . . .	5
2.2 HTCondor . . . . .	5
2.3 Open Science Grid . . . . .	6
<b>3 Related Work</b>	<b>8</b>
3.1 Comparison of existing mechanisms . . . . .	8
3.1.1 Condor Flocking . . . . .	8
3.1.2 Condor Job Router . . . . .	10
3.1.3 Co-Scheduler Solution . . . . .	12
3.2 Turnaround time based scheduling, IIT Roorkee . . . . .	14

<b>4</b>	<b>Design and Implementation of Co-Scheduler</b>	<b>15</b>
4.1	Capacity based scheduling . . . . .	17
4.2	Multi-site load distribution based scheduling . . . . .	18
4.3	Programming APIs . . . . .	19
4.3.1	Condor Log Reader and User API . . . . .	19
4.3.2	Synchronization Co-Scheduler Code . . . . .	21
<b>5</b>	<b>Evaluation</b>	<b>22</b>
<b>6</b>	<b>Conclusion</b>	<b>23</b>
	<b>Bibliography</b>	<b>24</b>

# List of Figures

1.1	Degradation occurrence: Co-Scheduling . . . . .	3
3.1	JobRouter: Transformation of Jobs . . . . .	10
4.1	Multi-site scheduling algorithm overview, classification of sites into slower and faster . . . . .	19

# List of Tables



# Chapter 1

## Introduction

Grid Computing as defined by wikipedia is the federation of computer resources from multiple locations to reach a common goal. The resources come in the form of hardware and software that allows us to submit jobs, run the jobs and monitor the jobs on the grid. Universities usually have multiple clusters across their campus and these are owned by different departments but stand united under the banner of university. Campus grids can be thought of as mini grids in which jobs are spanned across multiple clusters based on the need of the user and available computing infrastructure across multiple clusters within the computing resources across university.

Modern schedulers used in clusters provide innumerable features for policy making, resource management and scheduling, the problem of cluster performance degradation that occurs when either of the resource is throttled is a problem that hasn't been addressed. The problem of performance degradation when many jobs are scheduled on single system either based on processor equivalence or based on number of processor slots. Some of these schedulers like maui (CITE) are smart enough to take into account contention of other resources like RAM but ultimately convert the 2D vector values of CPU and RAM into single scalar value which equals to hard-coding the value or presenting these

resources in a ratio which makes us question effectiveness of such scheduling mechanism.

Here at Holland Computing Center, at University Of Nebraska-Lincoln we're tackling this issue of cluster degradation caused by over exploitation of one or more resource and I've come up with a solution by adaptively scheduling such highly dynamic resource on the grid across multiple sites by adaptively scaling with respect to degradation that is encountered and scheduling based on least turnaround time of the sites which would help increase the throughput of the overall Co-Scheduler which also loadbalances.

Existing schedulers depend on the availability of resources and frequent polling of it to determine the slots of scheduling. It should be noted that state of the art schedulers like maui/torque, slurm, condor take into account only CPU as a resource and the resources like RAM, I/O or Network are either ignored or their resource equivalent is converted into a scalar values which isn't an effective way of tackling the multiple resource scheduling problem as this might result in scheduling excessive jobs on single machine in case of schedulers ignoring the resource and loss of effective resources when schedulers convert the resource allocation into scalar values which means for example it'll be pre-assigned on every node that 1 CPU-CORE will have 2 GB of memory this kind of pre-assignment seems like its addressing the resource degradation of RAM but suffers from the above mentioned problems. We take a turnaround time approach to measure degradation, we define a resource to be degraded if and when we submit jobs that results in increasing the turnaround time by 25% . It must be noted that there isn't explicit measurement of individual resources of RAM, I/O, Network and CPU, there are no resource managers keeping track of these resource allocations. The Co-Scheduler is driven by the fact that turnaround time increases when concurrent jobs accessing these resources reaches a threshold value which in-turn causes degradation and this is the basis for our scheduling.

One aspect to Co-Scheduler is degradation detection and management by adapting to the throttling of resource another aspect is to efficiently distribute jobs across multiple sites

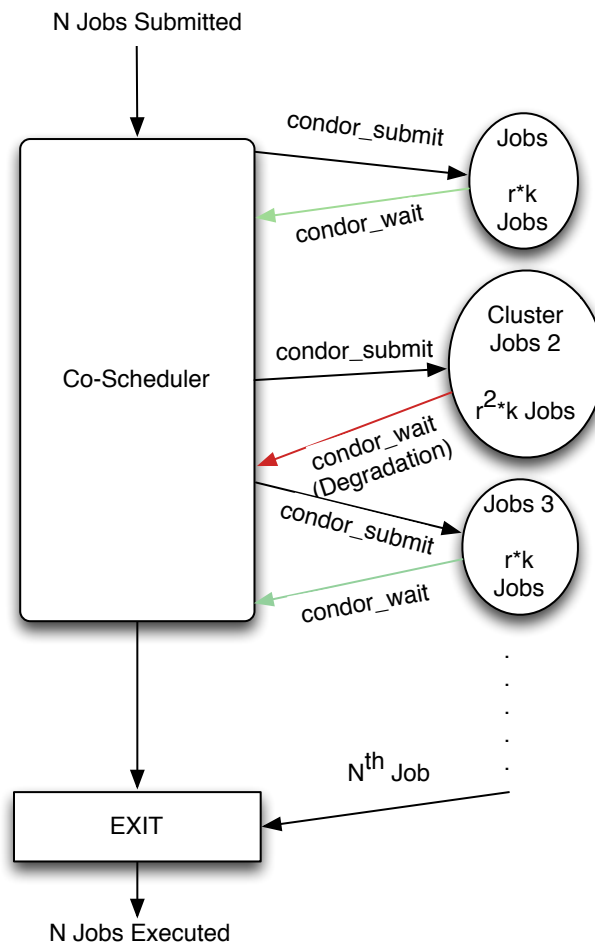


Figure 1.1: Degradation occurrence: Co-Scheduling

which increases the throughput of the workload. The Co-Scheduler adapts to degradation by backing off submission rates and waiting for random period of time for the next incremental submission and efficiently distributes the jobs and load-balances across multiple sites, thus submitting more jobs to a site with lesser turnaround time and also ensuring to submit lesser jobs to the sites with larger turnaround time thus efficiently load-balancing across multiple sites on a grid and improving throughput.

Finally, Grid environment is a heterogeneous environment, its absolutely necessary to use API's that are available on all the systems across the grid. To implement such a

Co-Scheduler we limit ourselves to condor based clusters and utilize libcondorapi. The result is a threaded Co-Scheduler that submits to multiple sites concurrently and is aware of multiple-resource degradation and loadbalances efficiently across multiple sites of the grid.

Please note that all the references to the Co-Scheduler in this thesis refers to the adaptive Co-Scheduler designed by me.

## **1.1 Co-Scheduler**

DHTC environment

# Chapter 2

## Background

### 2.1 High-Throughput Computing

High Throughput Computing, HTC is defined as a computing environment that delivers large amounts of computational power over a long period of time. The important factor being over a long period of time which differentiates HTC from HPC which focuses on getting large amount of work done in small amount of time. The workloads that run on condor system doesn't have an objective of how fast the job can be completed but how many times can the job be run in the next few months. In another definition of HTC, European Grid Infrastructure defines HTC as a computing paradigm that focuses on the efficient execution of large number of loosely coupled tasks.

### 2.2 HTCondor

HTCondor is a distributed system developed by HTCondor team at the University of Wisconsin-Madison. It provides High-Throughput Computing environment to sites that foster research computing and enables sites to share computing resources when comput-

ers are idle at a given site. HTCondor system includes a batch queuing system, scheduling policy, priority scheme, and resource classifications for a pool of computers mainly used for compute-intensive jobs, HTCondor runs on both UNIX and windows based workstations that are all connected by a network. Although there are other batch schedulers out there for dedicated machines. The power of condor comes from the fact that the amount of compute power represented by sum total of all the non-dedicated desktop workstations sitting on people's desks is sometimes far greater than the compute power of dedicated central resource. There are many unique tools and capabilities in HTCondor which make utilizing resources from non-dedicated systems effective. These capabilities include process checkpoint and migration, remote system calls and ClassAds. HTCondor suit also includes a powerful resource manager with an efficient match-making mechanism that is implemented via ClassAds, which makes HTCondor lucid when compared with other compute schedulers.

## 2.3 Open Science Grid

Open Science Grid(OSG), provides service and support for resource providers and scientific institutions using a distributed fabric of high throughout computational services. OSG was created to facilitate data analysis from the Large Hadron Collider . OSG doesn't own resources but provides software and services to users and enables opportunistic usage and sharing of resources among resource providers. The main goal of OSG is to advance science through open distributed computing. The OSG provides multi-disciplinary partnership to federate local, regional, community and national cyber-infrastructures to meet the needs of research and academic communities at all scales.

OSG provides resources and directions to Virtual Organizations(VO's) for the purposes of LHC experiments and HTC in general.

Building a OSG site requires listing background and careful planning. The major components of a OSG site includes a Storage Element and Compute Element.

Storage elements (SE) manage physical systems, disk caches and hierarchical mass storage systems, its an interface for grid jobs to underlying storage Storage Resource Management protocol and Globus Grid FTP protocol and others, A storage element requires an underlying storage system like hadoop, xrootd and a GridFTP server and an SRM interface.

A Compute Element(CE) allows grid users to run jobs on your site. It provides a bunch of services when run on the gatekeeper. The basic components include the GRAM and GridFTP on the same CE host to successfully enable file transfer mechanisms of Condor-G.

# Chapter 3

## Related Work

### 3.1 Comparison of existing mechanisms

There might be situations where we can have multiple condor pools and some of the pools might have many idle slots that are available for utilization. To efficiently utilize the resources across pools condor provides mechanisms like condor flocking and condor job router. These mechanisms provide similar functionality to the Co-Scheduler that I designed. In the following sections, we compare and evaluate these mechanisms with the design and functionality of Co-Scheduler.

#### 3.1.1 Condor Flocking

Flocking refers to a mechanism where jobs that cannot run in its own condor pool due to lack of resources runs in another condor pool where resources are available. Condor flocking enables load sharing between pools of computers. As pointed out by Campus Grids Thesis[CITE 3.1.2] flocking helps balance large workflows across different pools but not necessarily the jobs across these pools because of the scavenging and greedy nature of the condor scheduler.



`condor_schedd` advertises that it has idle jobs to the remote `condor_collector` and during the next phase of negotiation, if its found that there are computers available then they are allotted to the jobs in the matchmaking phase and the jobs then run on the remote pools. It so appears and the local job queue maintains as if the jobs are running locally.

Although Condor flocking has workflow balancing features across multiple condor pools it isn't aware of the slower and faster sites/pools. The adaptive Co-Scheduler keeps tabs on turnaround time and is aware of which site is faster or slower. If we look at the idea that scavenging idle resources increases throughput, it does but If we look at the case where jobs are greater than the available slots across all the pools, condor flocking stops working here as deeper understanding of sites would be required here to push more jobs at faster sites and importantly push less jobs at slower sites, its role is to scavenge idle computing slots across multiple pools.

Another aspect where condor flocking isn't designed to perform is when jobs are contending for the same resource (CPU, RAM, I/O & Network). Even though we'll be having idle slots on remote pools, condor flocking might not be able to use those slots as they'd be degraded because of the contention. In this case again Co-Scheduler comes in handy and based on the turnaround time waits for random amount of time for degradation to clear and doesn't put heavy job load on the degraded resource and diverts the jobs elsewhere which in-turn increases the throughput.

To conclude, we can say that condor flocking provides features of balancing large workflows and doesn't include features that would detect degradation in the cluster and also doesn't keep track of information of faster and slower sites, which might be exploited to increase the overall throughput of the system.

### 3.1.2 Condor Job Router

Condor manual defines the functions of job router to be the following:

The Condor Job Router is an add-on to the `condor_schedd` that transforms jobs from one type into another according to a configurable policy. This process of transforming the jobs is called job routing.

Condor Job Router can transform vanilla universe jobs to grid universe jobs and as it submits to multiple sites, the rate at which it starts submitting equals rate at which the sites execute them thus providing platform to balance large workflows across multiple grid sites and replenishing the jobs at faster sites once they get done. Job router sends more jobs to a site if the jobs submitted are not idle and stops submitting jobs if the submitted jobs sit idle on the remote cluster and Job router is not aware about which site is faster or slower.

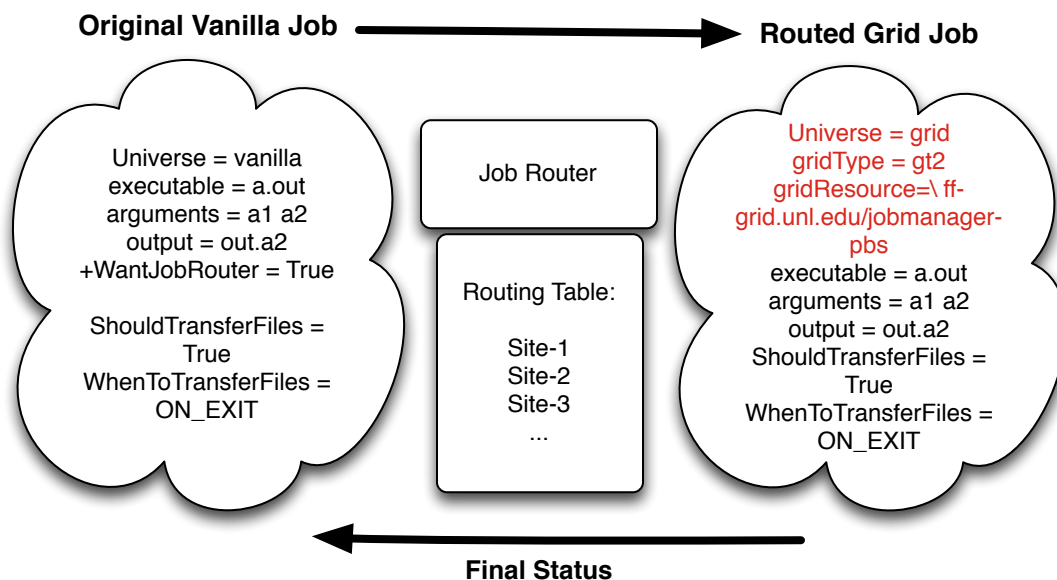


Figure 3.1: JobRouter: Transformation of Jobs

A job is transformed to the grid universe by making a copy of the original job ClassAd, modifying some attributes of the job , this copy is called the routed copy and this routed copy shows up in the job queue with a new job id.

Condor job router utilizes routing table which contains the listings of sites the job must be submitted to and the name of the grid resource declared and processed during condor config file which is defined by the new ClassAds.

```

1  # Now we define each of the routes to send jobs on
2  JOB_ROUTER_ENTRIES = \
3      [ GridResource = "gt5 ff-grid.unl.edu/jobmanager-pbs"; \
4        name = "Firefly"; \
5      ] \
6      [ GridResource = "gt5 tusker-gw1.unl.edu/jobmanager-pbs"; \
7        name = "Tusker"; \
8      ] \
9      [ GridResource = "gt5 pf-grid.unl.edu/jobmanager-condor"; \
10       name = "Prairiefire"; \
11     ]\

```

Condor Job Router seems to be a step up from the Condor Flocking in terms of scavenging resources and sending the extra jobs to another condor pool. Condor Job Router also maintains the rate of the jobs on submit hosts equal to that of remote clusters. But the job router does not keep track of how fast condor cluster is. Thus condor job router does not optimize the displacement of jobs to slower cluster, since it maintains the rate of submission of jobs, we can be sure that if a job is complete at a faster site, it is immediately replaced by the next one. The same is true for a job at slower site too which might result in increasing the degradation if there exists some, for example if the job

router kept track of slower sites, it can send less jobs to slower sites there by increasing the overall throughput of the given workflow.

A preliminary examination shows that condor job router does seem to have degradation detection features, it does not submit to a pool that already has idle jobs but a close examination reveals that even though it submits to a pool with non-idle jobs it isn't aware of the fact that the jobs in the queue would have undergone degradation due to contention on same resource and thus isn't a viable mechanism for degradation detection.

### 3.1.3 Co-Scheduler Solution

Due to the heterogeneous nature of the grid there could be problems with resource sharing on the grid I came up with a solution that primarily solves the problem of degradation due to contention among jobs for any given resource like RAM,I/O and network. The contention among jobs sometimes give rise to degradation depending upon the availability of the resource that results in reduced performance of the cluster or crashes the cluster altogether.Efficient and high throughput distribution of jobs across multiple cluster is another challenging problem that we have solved in this particular Co-Scheduler Solution.

In the first part of the solution we target the basic problem of degradation to give an example, suppose we have a file server that can serve 100 MBps of data which is available for concurrent access and user has submitted 100 jobs each consuming 10 MBps of data I/O if 100 slots of CPU slots are available a typical state of the art Scheduler schedules these 200 jobs without actually looking at the I/O load, this results in a degraded system, overtime if more jobs are Scheduled to utilize this file server, this file serve might even crash.We would need a degradation handling mechanism that would adaptively scale with the load and exponentially backoff during high contention period thus we need an

intervention in the form of Co-Scheduler that intelligently handles degradation. We start the capacity algorithm by submitting jobs to the cluster and measure the turnaround time of each iteration if we find that the current turnaround time of an iteration is 25% greater than the previous iteration we term it as a degradation. We now try to find the best capacity of jobs between the current and the previous iteration by exponentially backing-off between these two iterations. We keep finding the optimal capacity dynamically for each iteration as the contention may change and optimal capacity keeps varying. Once the optimal capacity is found then it is retained for job submission for random amount of time limited to maximum of ten iterations of jobs thus solving the problem of performance degradation.

The other problem that we tackled was that of efficient distribution of work load to multiple sites of the grid, some of these sites could be error prone, some of them could be faster some of them can be slower, we need to take an approach that improves the overall throughput of the workload, so we keep track of average turnaround time for each batch of jobs submitted to the cluster. We take average of turnaround time of the batch of job that is submitted across multiple cluster we group turnaround times of these batches of jobs based on the value of the turnaround time to be greater than, lesser than or equal to the average turnaround time of all batch of jobs. This gives us a classification of sites that run faster, we exploit this information in our scheduler. To implement multi-site jobs submission and load balancing feature we start submitting one job to each cluster and measure the turnaround time once the job returns. Now we have the information of turnaround time of a particular resource of all the clusters, although continuing to submit jobs concurrently in a similar manner would result in the completion of the workload in a non efficient manner, out of the list of all the turnaround time of different sites we take the sites that have lower than the average turnaround time of all batches of jobs and increase the job propagation factor to these jobs which in-turn increases the number

of jobs submitted to that site, which helps us efficiently distribute workload to faster clusters than the slower clusters ( submitting to faster site would enable us to submit more number of jobs to the particular site. As the turnaround time is lower we would finish jobs quickly and make room for more number of jobs if we submit less to slower sites that can in turn help us improve the throughput to a greater extent by enabling us to stop further degradation at those sites and help us submit more to faster sites . ) Error handling mechanism is gracefully handled by condor and by this approach we have an efficient system with increased throughput and with proper distribution of load across multiple clusters.

### **3.2 Turnaround time based scheduling, IIT Roorkee**

## Chapter 4

# Design and Implementation of Co-Scheduler

The co-scheduler requires the presence of a condor installation and libcondorapi. It is written in the C++ language and has extensively made use of pthreads for synchronization and multithreading. The co-scheduler has two components to it, the first one is capacity detection algorithm which is found by measuring degradation and other aspect to it is multi-site workload distribution across grid. The following paragraphs detail the engineering aspect of the design. The program begins by taking number of jobs, sites information and submit script information as its input. These number of jobs are the ones submitted across multiple sites. The next input file contains information on the list of sites that can be used for load balancing in the GridResource format of the condor ClassAds API, like *tusker-gw1.unl.edu/jobmanager-pbs*. It is assumed that all the sites listed in the sites input file are working and do not have any misconfiguration issues. The third and the final option to the scheduler is the submit description file. All the job ClassAd information and requirements can be written in this section and all of these will be applicable on the grid when a particular job is scheduled. The two main important

attributes required in the submit description file is *universe* which needs to be grid, all the time as we are submitting to grid sites and the other most important thing is the *grid-proxy*. We need a proxy to delegate our identification on the grid. *voms-proxy-init* is used to generate voms proxy at the submit host, we can specify `--voms` as our virtual organization, `hcc:/hcc` in this case and `--hours` as the number of hours the proxy would be active/valid, it must be ensured that the proxy period is approximately greater than the length of period of runtime of jobs for smooth running of jobs.

HTCondor-G is the grid environment, when jobs are sent across grid universe using Globus software. Globus toolkit provides support for building grid systems. Submitting, managing and executing jobs have same capabilities in both HTCondor and HTCondor-G world. Globus provides fault tolerant features to HTCondor-G jobs. GRAM is Grid Resource Allocation and Management protocol, supports remote submission of computational request.

`gt2` is an initial GRAM protocol which is used in Globus Toolkit version 1 and 2. `gt2` is also referred to as the pre-web services GRAM or GRAM2.

`gt5` is the latest GRAM protocol, which is an extension of GRAM2 and is intended to be more scalable and robust, referred to as GRAM5.

As we know that `condor_wait` can be used to wait for a certain job or number of jobs to complete, it watches the log file and sees if the completion entry of the job is made in the log file that is generated by the `log` command in the condor submit description file. There are two options we can specify to `condor_wait` one is `-num`, number-of-jobs that waits till the number-of-jobs are completed and the other option is to specify `-wait`, seconds that waits for seconds amount of time, if nothing is specified `condor_wait` waits indefinitely till the job(s) are complete.



## 4.1 Capacity based scheduling

This algorithm detects degradation and once we detect degradation we find the optimal capacity:

---

**Algorithm 1** Algorithm for determining optimal capacity by detecting degradation

---

```

c1  $\leftarrow$  1.5 {c1: Job propagation factor}
c2  $\leftarrow$  0.25 {c2: Degradation factor}
while true do
  if  $T_2 < c_2 * T_1$  then
    jobSubmission(k)
     $T_1 = (T_1 + T_2) / 2$ 
  end if
  if  $T_2 > c_2 * T_1$  then
    degradation_high  $\leftarrow$  k
    degradation_low  $\leftarrow$  k/2
    optimalCapacity(degradation_high, degradation_low)
  end if
end while

```

---



---

**Algorithm 2** Algorithm for determining optimal capacity by detecting degradation

---

```

mid  $\leftarrow$  (high + low) / 2
k  $\leftarrow$  mid
jobSubmission(k)
if  $T_2 < c_2 * T_1$  then
   $T_1 \leftarrow (T_1 + T_2) / 2$ 
  optimalCapacity(mid, high);
end if
if  $T_2 > c_2 * T_1$  then
  optimalCapacity(low, mid);
end if
return mid

```

---

Diagram:

## 4.2 Multi-site load distribution based scheduling

Multi-site load distribution based scheduling uses the turnaround time at each site for sending future jobs, its a threaded algorithm, the following is extracted function per thread:

---

**Algorithm 3** Algorithm for distribution of workflow load across multiple sites on the grid

---

```

high  $\leftarrow$  0
low  $\leftarrow$  0
average  $\leftarrow$  0
{arrayMultipleSitesTime is populated with turnaround times of all sites}
arrayMultipleSitesTime.sort()
low = arrayMultipleSitesTime.front()
high = arrayMultipleSitesTime.back()
average = sum of all elements of arrayMultipleSitesTime / Size of arrayMultipleSites-
Time
if turnaroundTime.Thread == low then
    c1  $\leftarrow$  c1 * 4
end if
if turnaroundTime.Thread == high then
    c1  $\leftarrow$  c1/2?(c1/2 : 1)
end if
if turnaroundTime.Thread < average then
    c1  $\leftarrow$  c1 * 2
end if

```

---

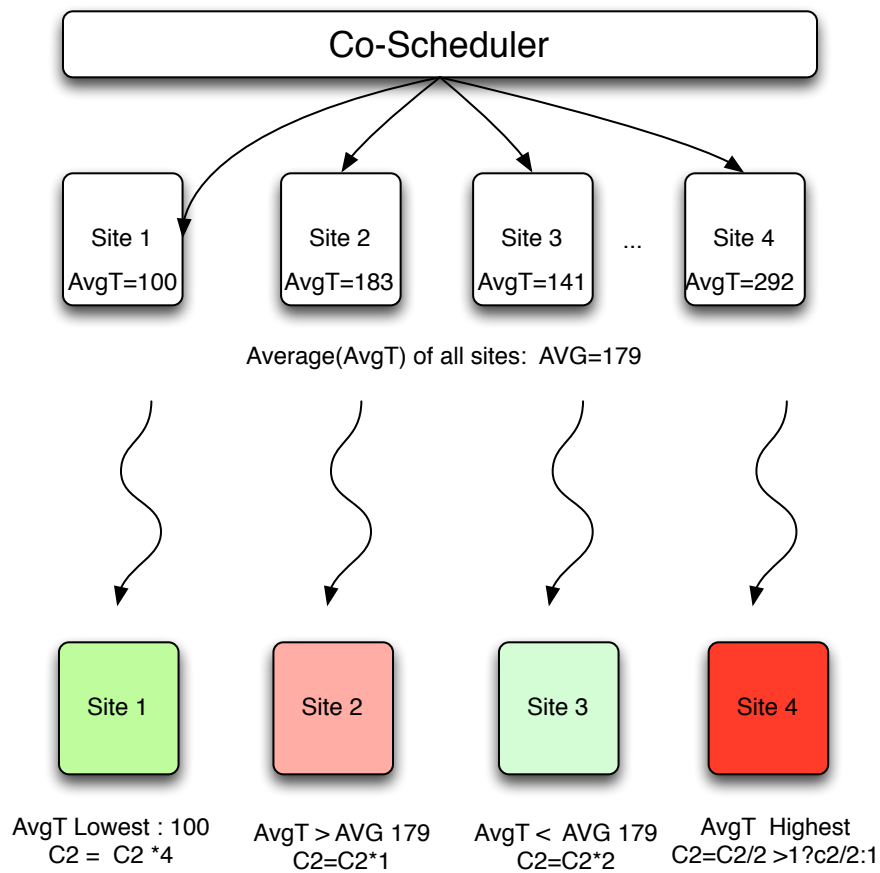


Figure 4.1: Multi-site scheduling algorithm overview, classification of sites into slower and faster

Diagram:

## 4.3 Programming APIs

### 4.3.1 Condor Log Reader and User API

Constructor: `ReadUserLog reader(fp,false,false);`

`ULogEventOutcome` (defined in `condor_event.h`):

Status events for job detection:

ULOG\_OK: Event is valid

ULOG\_NO\_EVENT: No event occurred (like EOF)

ULOG\_RD\_ERROR: Error reading log file

ULOG\_MISSED\_EVENT: Missed event

ULOG\_UNK\_ERROR: Unknown Error

Job Submission Algorithm:

```

1 void wait_calculate(string hostFile, args *data, int nSites) {
2     FILE *fp;
3     ReadUserLog reader(fp,false,false);
4     ULogEvent *event = NULL;
5     if(reader.isInitialized()) {
6         cout<<"Initialized reading log file\n"<<flush;
7     }
8     else{
9         cout<<"ERROR! reading log file host.log"<<flush;
10    }
11
12    while(reader.readEvent(event)==ULOG_OK) {
13        ostringstream cluster_proc;
14        if((*event).eventNumber==ULOG_EXECUTE ) {
15            ExecuteEvent *exec
16            = static_cast<ExecuteEvent*>(event);
17            count++;

```

```

18         //condor_wait -num K, where K is the
19         //amount of jobs completed till the wait.
20         if(count<=N) {
21             char tmp[100];
22             sprintf(tmp,"condor_wait -num
23             \"%d %s\",count,hostFile.c_str());
24             return_value=system(tmp);
25         }
26         continue;
27     }
28     if((*event).eventNumber==ULOG_JOB_TERMINATED) {
29         JobTerminatedEvent *term
30         = static_cast<JobTerminatedEvent*>(event);
31
32         if(term->normal) {
33             //Works only for local jobs,
34             //find the CPU time of local jobs
35         }
36     }
37 }
38
39 fclose(fp);
40 }

```

### 4.3.2 Synchronization Co-Scheduler Code

# **Chapter 5**

## **Evaluation**

## **Chapter 6**

## **Conclusion**

# Bibliography

- [1] Condor manual, section 5.6.
- [2] D.H.J. Epema, M. Livny, R. van Dantzig, X. Evers, and J. Pruyne. A worldwide flock of Condors: Load sharing among workstation clusters. *Future Generation Computer Systems*, 12(1):53–65, 1996.
- [3] I. Foster and C. Kesselman. Globus: A metacomputing infrastructure toolkit. *International Journal of High Performance Computing Applications*, 11(2):115, 1997.
- [4] J. Frey, T. Tannenbaum, M. Livny, I. Foster, and S. Tuecke. Condor-G: A computation management agent for multi-institutional grids. *Cluster Computing*, 5(3):237–246, 2002.
- [5] Cluster Resources Inc. Cluster resources :: Products - TORQUE Resource Manager:, January 2011. <http://www.clusterresources.com/pages/products/torque-resource-manager.php>.
- [6] Miron Livny and Rajesh Raman. High-throughput resource management. In Ian Foster and Carl Kesselman, editors, *The Grid: Blueprint for a New Computing Infrastructure*. Morgan Kaufmann, 1998.
- [7] LLNL. Slurm reference.



- [8] Holland Computing Center. Holland computing center, January 2011. <http://hcc.unl.edu/>.
- [9] R. Pordes, D. Petravick, B. Kramer, D. Olson, M. Livny, A. Roy, P. Avery, K. Blackburn, T. Wenaus, et al. The Open Science Grid. In *Journal of Physics: Conference Series*, volume 78, page 012057. IOP Publishing, 2007.
- [10] I. Sfiligoi. glideinWMS—a generic pilot-based workload management system. In *Journal of Physics: Conference Series*, volume 119, page 062044. IOP Publishing, 2008.
- [11] D. Thain, T. Tannenbaum, and M. Livny. Distributed computing in practice: The Condor experience. *Concurrency and Computation: Practice and Experience*, 17(2-4):323–356, 2005.
- [12] M. Zvada, D. Benjamin, and I. Sfiligoi. CDF GlideinWMS usage in Grid computing of high energy physics. In *Journal of Physics: Conference Series*, volume 219, page 062031. IOP Publishing, 2010.